

Dave Gabriel Jazmin
Junmar Perez
Aaron Bonaobra

LAB 3 DOCUMENTATION

1.1 Objective

The purpose of this documentation is to outline the steps taken to preprocess a dataset using Apache Spark. This includes data loading, cleaning, transformation, and aggregation to derive meaningful insights.

1.2 Introduction

Apache Spark is a powerful distributed computing system used for big data processing. It provides efficient ways to handle large datasets using the PySpark library in Python.

2.1 Dataset Information

The dataset used in this project contains information about car prices and specifications, including brand, model, year, engine size, fuel type, transmission, mileage, doors, owner count, and price.

3.1 Loading the Dataset

The dataset was loaded into a Spark DataFrame using the following command:

```
1. from pyspark.sql import SparkSession
2. from pyspark.sql.functions import col, avg, regexp_replace, when
3. import os
4. import shutil
5.
6. spark = SparkSession.builder.appName("CarDataProcessing").getOrCreate()
7.
8. file_path = "C:\\Users\\daveg\\Desktop\\SKWELA\\3rd Year (2nd Sem)\\Big
   Data\\car_price_dataset.csv"
9. df = spark.read.csv(file_path, header=True, inferSchema=True)
10.
11. print("Original Dataset:")
12. df.show(5)
```

```

>>> from pyspark.sql import SparkSession
>>> from pyspark.sql.functions import col, avg, regexp_replace, when
>>> import os
>>> import shutil
>>>
>>> spark = SparkSession.builder.appName("CarDataProcessing").getOrCreate()
25/02/17 23:27:06 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.
>>>
>>> file_path = r"C:\Users\daveg\Desktop\SKWELA\3rd Year (2nd Sem)\Big Data\car_price_dataset.csv"
>>> df = spark.read.csv(file_path, header=True, inferSchema=True)
>>>
>>> print("Original Dataset:")
Original Dataset:
>>> df.show(5)
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Brand | Model | Year | Engine_Size | Fuel_Type | Transmission | Mileage | Doors | Owner_Count | Price |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Kia   | Rio   | 2020 | 4.2         | Diesel    | Manual       | 289944  | 3     | 5           | 8501  |
| Chevrolet | Malibu | 2012 | 2.0         | Hybrid    | Automatic    | 5356    | 2     | 3           | 12092 |
| Mercedes | GLA   | 2020 | 4.2         | Diesel    | Automatic    | 231440  | 4     | 2           | 11171 |
| Audi  | Q5    | 2023 | 2.0         | Electric  | Manual       | 160971  | 2     | 1           | 11780 |
| Volkswagen | Golf | 2003 | 2.6         | Hybrid    | Semi-Automatic | 286618 | 3     | 3           | 2867  |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```

3.2 Data Cleaning

Several preprocessing steps were performed:

- Removing duplicate records
 - Handling missing and incorrect values
 - Converting data types for consistency
13. `df_cleaned = df.dropDuplicates()`
 14. `df_cleaned = df_cleaned.withColumn("price", regexp_replace(col("price"), "[^0-9.]", ""))`
 15. `df_cleaned = df_cleaned.withColumn("year", regexp_replace(col("year"), "[^0-9]", ""))`
 16. `df_cleaned = df_cleaned.withColumn("price", when(col("price") == "", "0").otherwise(col("price")))`
 17. `df_cleaned = df_cleaned.withColumn("year", when(col("year") == "", "0").otherwise(col("year")))`
 - 18.
 19. `df_cleaned = df_cleaned.fillna({"price": "0", "brand": "Unknown", "year": "0"})`
 20. `df_cleaned = df_cleaned.withColumn("price", col("price").cast("float"))`
 21. `df_cleaned = df_cleaned.withColumn("year", col("year").cast("int"))`

3.3 Filtering Data

A filter was applied to remove records where the price was below 5000.

22. `df_filtered = df_cleaned.filter(col("price") > 5000)`
23. `print("Cleaned and Filtered Dataset:")`
24. `df_filtered.show(5)`

```
>>> df_filtered = df_cleaned.filter(col("price") > 5000)
>>>
>>> print("Cleaned and Filtered Dataset:")
Cleaned and Filtered Dataset:
>>> df_filtered.show(5)
```

Brand	Model	year	Engine_Size	Fuel_Type	Transmission	Mileage	Doors	Owner_Count	price
Chevrolet	Impala	2013	4.8	Petrol	Manual	206381	4	4	8672.0
Toyota	Camry	2023	1.5	Diesel	Manual	117049	2	5	10159.0
Audi	A3	2005	2.3	Electric	Manual	112828	3	2	7643.0
Volkswagen	Passat	2007	3.0	Electric	Manual	272439	3	3	5751.0
Volkswagen	Golf	2010	2.9	Electric	Manual	249910	5	3	7001.0

only showing top 5 rows

3.4 Aggregation

The average price of cars was calculated based on the brand.

```
25. df_grouped = df_filtered.groupBy("brand").agg(avg("price").alias("Average_Price"))
26. print("Average Price by Brand:")
27. df_grouped.show()
```

3.5 Saving the Processed Data

The cleaned and aggregated dataset was saved as a CSV file.

```
28. output_path = "C:\\Users\\daveg\\Desktop\\SKWELA\\3rd Year (2nd Sem)\\Big
    Data\\car_price_dataset"
29. df_grouped.write.csv(output_path, header=True)
30. print(f'Saved file path: {output_path}')
```

```
>>> df_grouped = df_filtered.groupBy("brand").agg(avg("price").alias("Average_Price"))
>>>
>>> print("Average Price by Brand:")
Average Price by Brand:
>>> df_grouped.show()
```

brand	Average_Price
Volkswagen	9610.589800443458
Kia	9483.450285714285
Chevrolet	9649.525139664804
Hyundai	9502.603686635945
Honda	9349.710585585586
Audi	9682.598233995584
Mercedes	9653.273053892215
BMW	9459.895402298851
Toyota	9585.8138424821
Ford	9528.752155172413

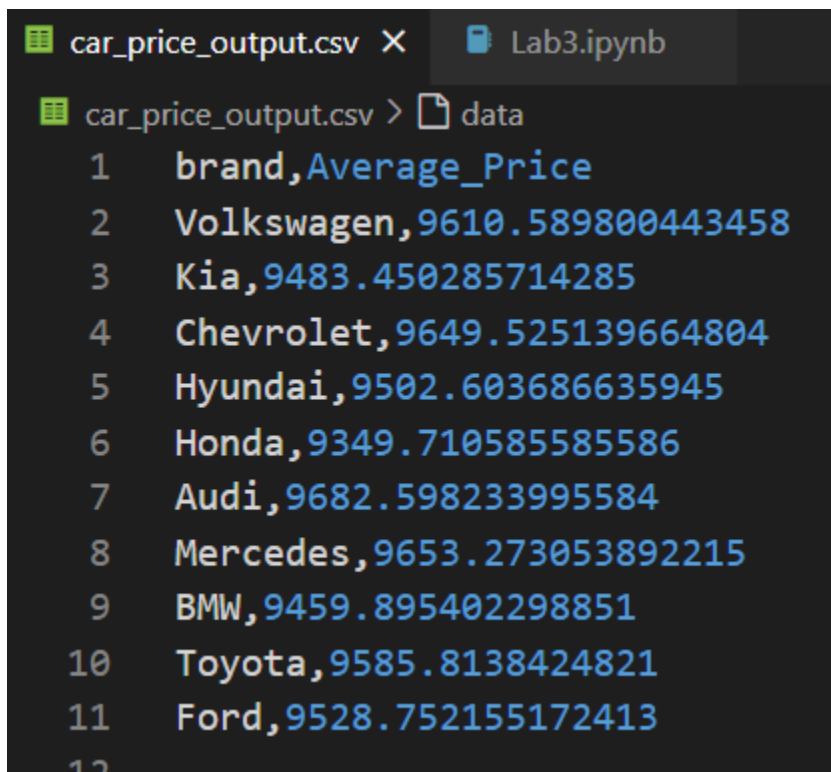
4. Results

The final output consists of a cleaned dataset and an aggregated report showing the average price of cars by brand.

```
output_path = r"C:\For School\Big Data\Lab3\car_price_output.csv"

# Convert to Pandas and save as a single CSV
df_grouped.toPandas().to_csv(output_path, index=False)

print(f"Saved file path: {output_path}")
```



```
car_price_output.csv X Lab3.ipynb

car_price_output.csv > data
1  brand,Average_Price
2  Volkswagen,9610.589800443458
3  Kia,9483.450285714285
4  Chevrolet,9649.525139664804
5  Hyundai,9502.603686635945
6  Honda,9349.710585585586
7  Audi,9682.598233995584
8  Mercedes,9653.273053892215
9  BMW,9459.895402298851
10 Toyota,9585.8138424821
11 Ford,9528.752155172413
12
```

5. Challenges & Solutions

Challenges:

1. **Unicode Error in File Path** - The error `unicodeescape` was encountered due to incorrect handling of backslashes in the file path.
2. **Data Type Mismatches** - Some numeric columns were read as strings, leading to type conversion errors.

6. Conclusion

In this project, we successfully loaded, cleaned, transformed, and analyzed a car dataset using Apache Spark. The data was processed efficiently, and the final aggregated results provided insights into average car prices by brand. This demonstrates the power of PySpark in handling large-scale data processing tasks.